

# Computer Networks - Assignment 2

*Nimitt (22110169)*

*Sumeet Sawale (22110234)*

## Task 1: Comparison of congestion control protocols

---

### Setup and Tools Used


- UTM Virtual Machine Ubuntu 20.04
- Mininet
- Wireshark

### Implementation

---

#### Setting Up the Network

- We use Mininet to setup the desired network configuration as per the input option:



```
def setup_topology(option, link_loss=None):
```

```
    class CustomTopo(Topo):
```

```
        def build(self):
```

```
            s1 = self.addSwitch('s1')
```

```
            s2 = self.addSwitch('s2')
```

```
            s3 = self.addSwitch('s3')
```

```
            s4 = self.addSwitch('s4')
```

```
            h1 = self.addHost('h1')
```

```
            h2 = self.addHost('h2')
```

```
            h3 = self.addHost('h3')
```

```
            h4 = self.addHost('h4')
```

```
            h5 = self.addHost('h5')
```

```
            h6 = self.addHost('h6')
```

```
            h7 = self.addHost('h7')
```

```
            self.addLink(h1, s1)
```

```
            self.addLink(h2, s1)
```

```
            self.addLink(h3, s2)
```

```

        self.addLink(h4, s2)
        self.addLink(h5, s3)
        self.addLink(h6, s3)
        self.addLink(h7, s4)

        self.addLink(s1, s2)
        self.addLink(s2, s3)
        self.addLink(s3, s4)

    net = Mininet(topo=CustomTopo(), controller=OVSController)
    net.start()

    for switch in ['s1', 's2', 's3', 's4']:
        net.get(switch).cmd('ovs-ofctl add-flow %s actions=normal' % switch)

    if option in ['c', 'd']:
        net.get('s1').cmd('tc qdisc add dev s1-eth1 root tbf rate 100Mbit burst 10kb l
        net.get('s2').cmd('tc qdisc add dev s2-eth2 root tbf rate 50Mbit burst 5kb l
        net.get('s3').cmd('tc qdisc add dev s3-eth3 root tbf rate 100Mbit burst 10kb

        if option == 'd' and link_loss is not None:
            net.get('s2').cmd(f'tc qdisc add dev s2-eth2 root netem loss {link_loss}%

    return net

```

## Running Test

- As per the option, we set up the configuration and save the network flow statistics for further metric computation:

```

def run_iperf_test(net, option, cc_scheme, condition=None, link_loss=None):
    server = net.get('h7')
    server.cmd('tcpdump -i h7-eth0 port 5001 -w iperf_capture.pcap &')
    server.cmd('iperf3 -s -p 5001 &')
    time.sleep(2)

```

```

clients = {
    'a': [('h1', 0, 150)],
    'b': [('h1', 0, 150), ('h3', 15, 120), ('h4', 30, 90)],
    'c': { # Map conditions to client-server setups
        'a': [('h1', 0, 150), ('h2', 0, 150)],
        'b': [('h1', 0, 150), ('h3', 0, 150)],
        'c': [('h1', 0, 150), ('h3', 0, 150), ('h4', 0, 150)]
    },
    'd': { # Loss configurations
        'a': [('h1', 0, 150), ('h2', 0, 150)],
        'b': [('h1', 0, 150), ('h3', 0, 150)],
        'c': [('h1', 0, 150), ('h3', 0, 150), ('h4', 0, 150)]
    }
}

test_clients = clients[option] if option in ['a', 'b'] else clients[option][condition]

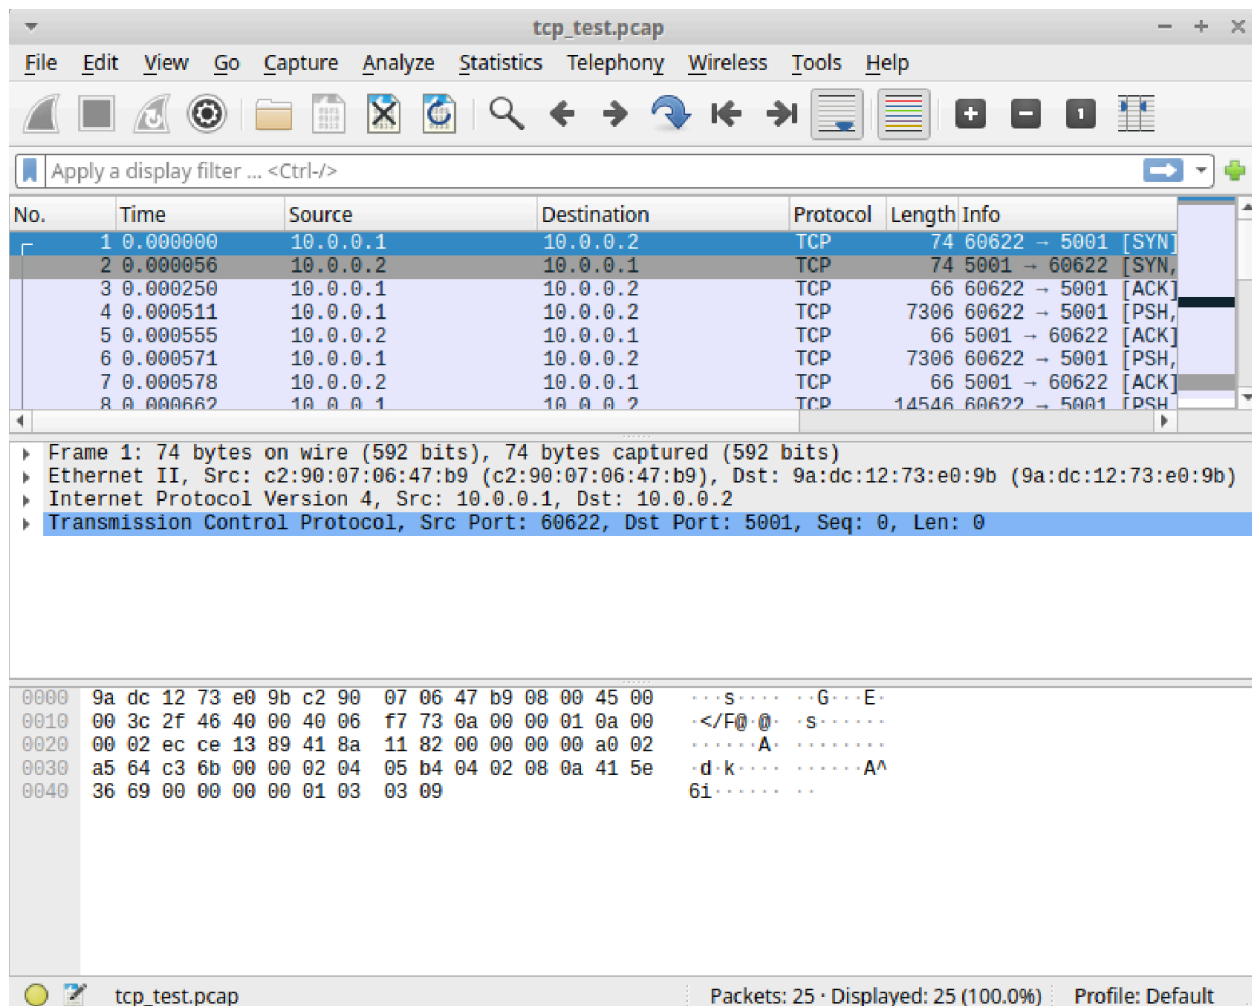
for host, delay, duration in test_clients:
    time.sleep(delay)
    client = net.get(host)
    client.cmd(f'iperf3 -c 10.0.0.7 -p 5001 -b 10M -P 10 -t {duration} -C {cc_scheme}')

time.sleep(160)
server.cmd('pkill tcpdump')
analyze_pcap()
print("Tests completed! Results saved to iperf_capture.pcap and *_{cc_scheme}")

```

## Analysing Metrics

- We analyse and compute the required metrics using saved network flow statistics using Wireshark GUI and tshark:



```
def analyze_pcap():
    print("Analyzing pcap file for performance metrics...")
    throughput = subprocess.getoutput("tshark -r iperf_capture.pcap -q -z io,stat,
    goodput = subprocess.getoutput("tshark -r iperf_capture.pcap -Y 'tcp.len > 0'
    packet_loss = subprocess.getoutput("tshark -r iperf_capture.pcap -q -z exper
    max_packet_size = subprocess.getoutput("tshark -r iperf_capture.pcap -T fiel

    print("Throughput:", throughput)
    print("Goodput (total data packets received):", goodput)
    print("Packet loss rate:", packet_loss)
    print("Maximum packet size achieved:", max_packet_size)
```

# Results

---

## Part (a)

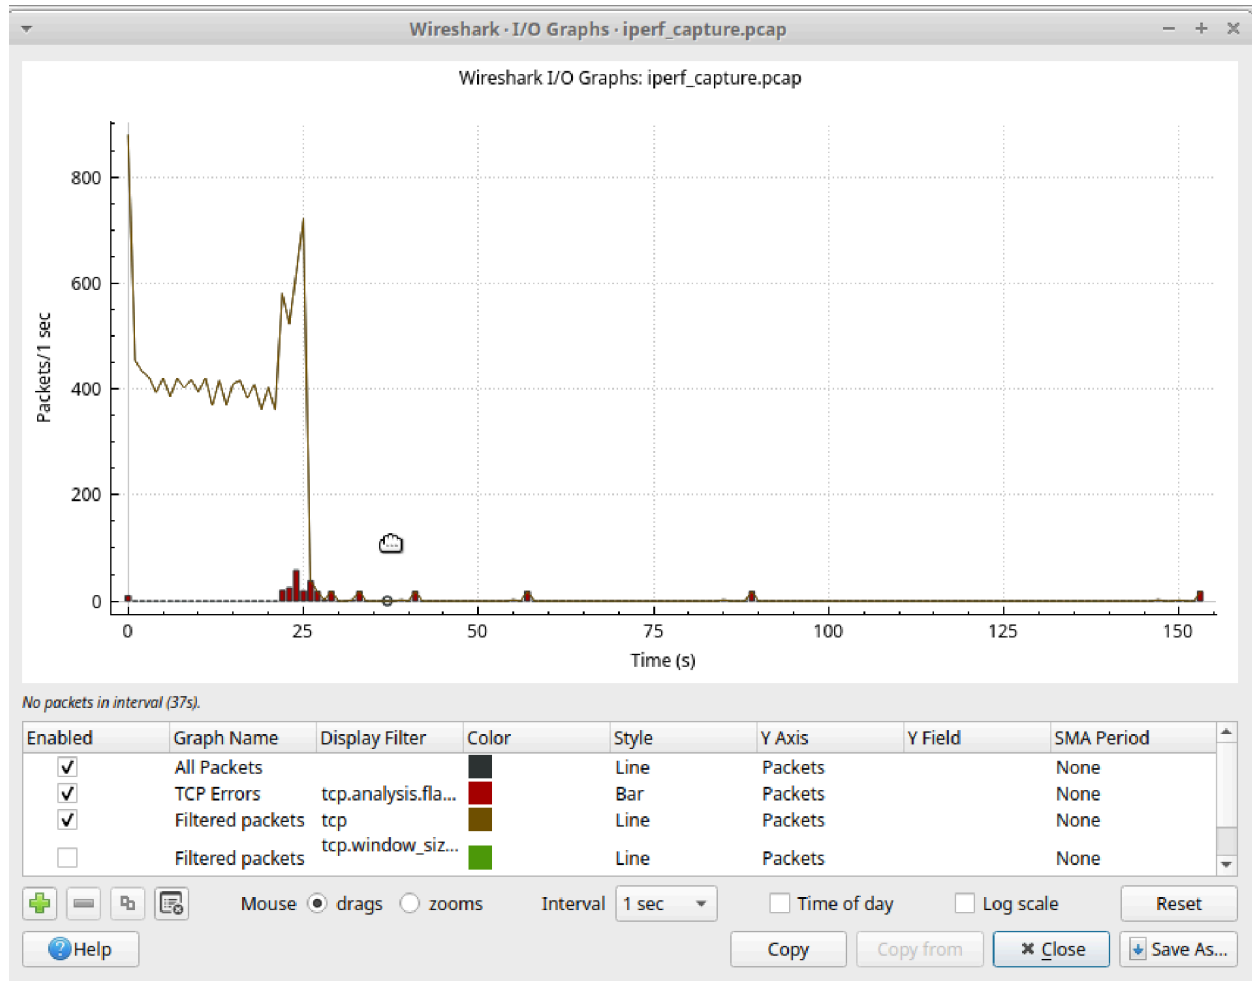
- Congestion Control = 'yeah'

Throughput:

Using Tshark

Time Interval (s)	Packets	Bytes
0 <> 10	4626	126135170
10 <> 20	3947	124778902
20 <> 30	3284	72659524
30 <> 40	22	1452
40 <> 50	18	1188
50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	20	1320
90 <> 100	0	0
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	2	132
150 <> Dur	20	1321

IO Graph



Goodput:

We define goodput as,

$$\frac{\text{Good Packets}}{\text{Total Packets}} * 100$$

= 67%

Packet Loss: 0%

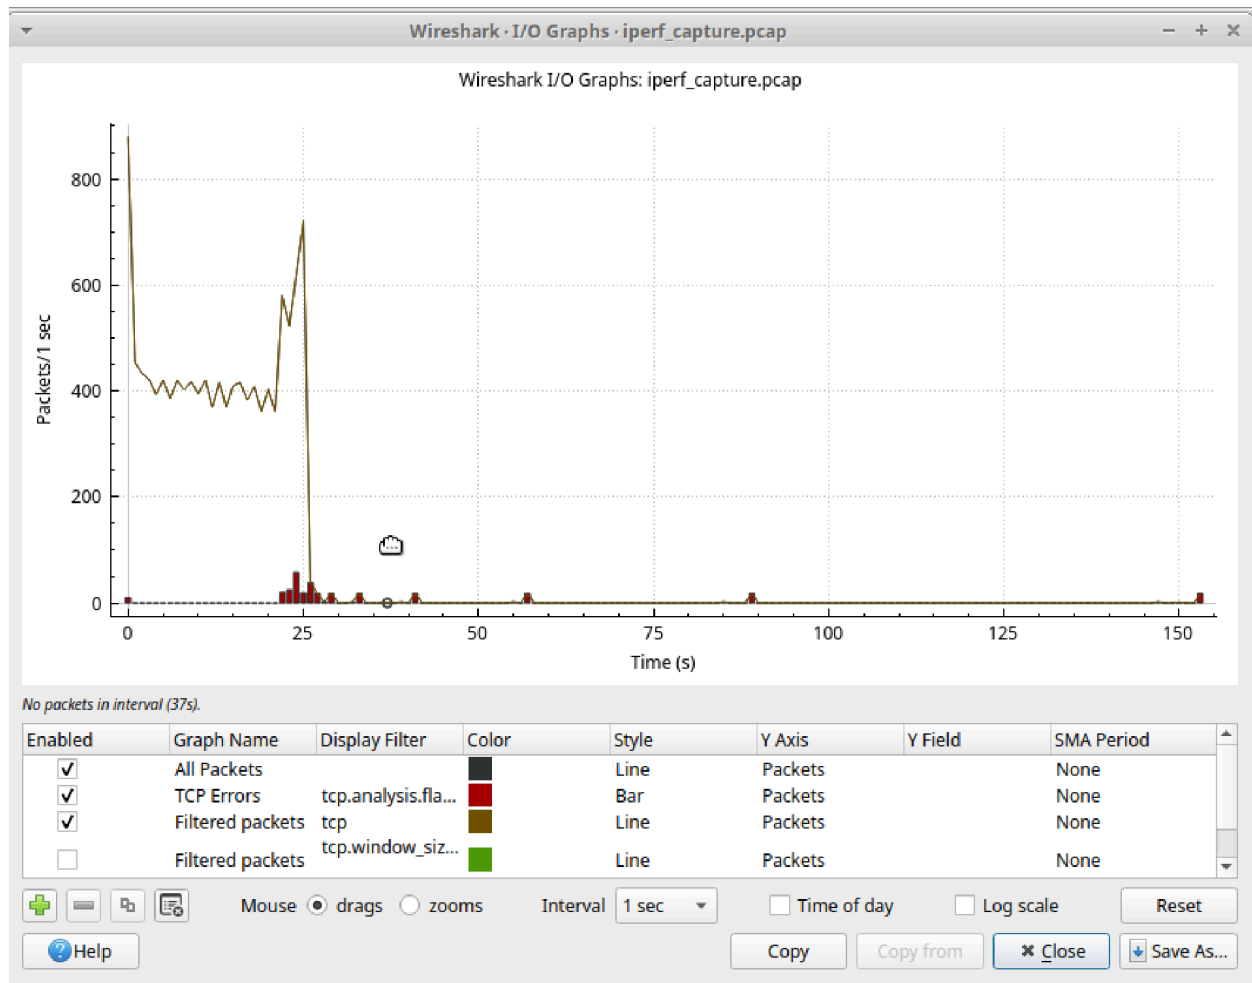
Maximum Packet Size: 65226 bytes

- Congestion Control = 'bbr'

Throughput:

Time Interval (s)	Packets	Bytes
0 <> 10	4821	126154475
10 <> 20	4330	124804180
20 <> 30	4168	73388664
30 <> 40	89	1764594
40 <> 50	14	924
50 <> 60	18	1188
60 <> 70	2	132
70 <> 80	0	0
80 <> 90	16	1056
90 <> 100	4	264
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	10	660
150 <> Dur	12	793





Goodput: 64%

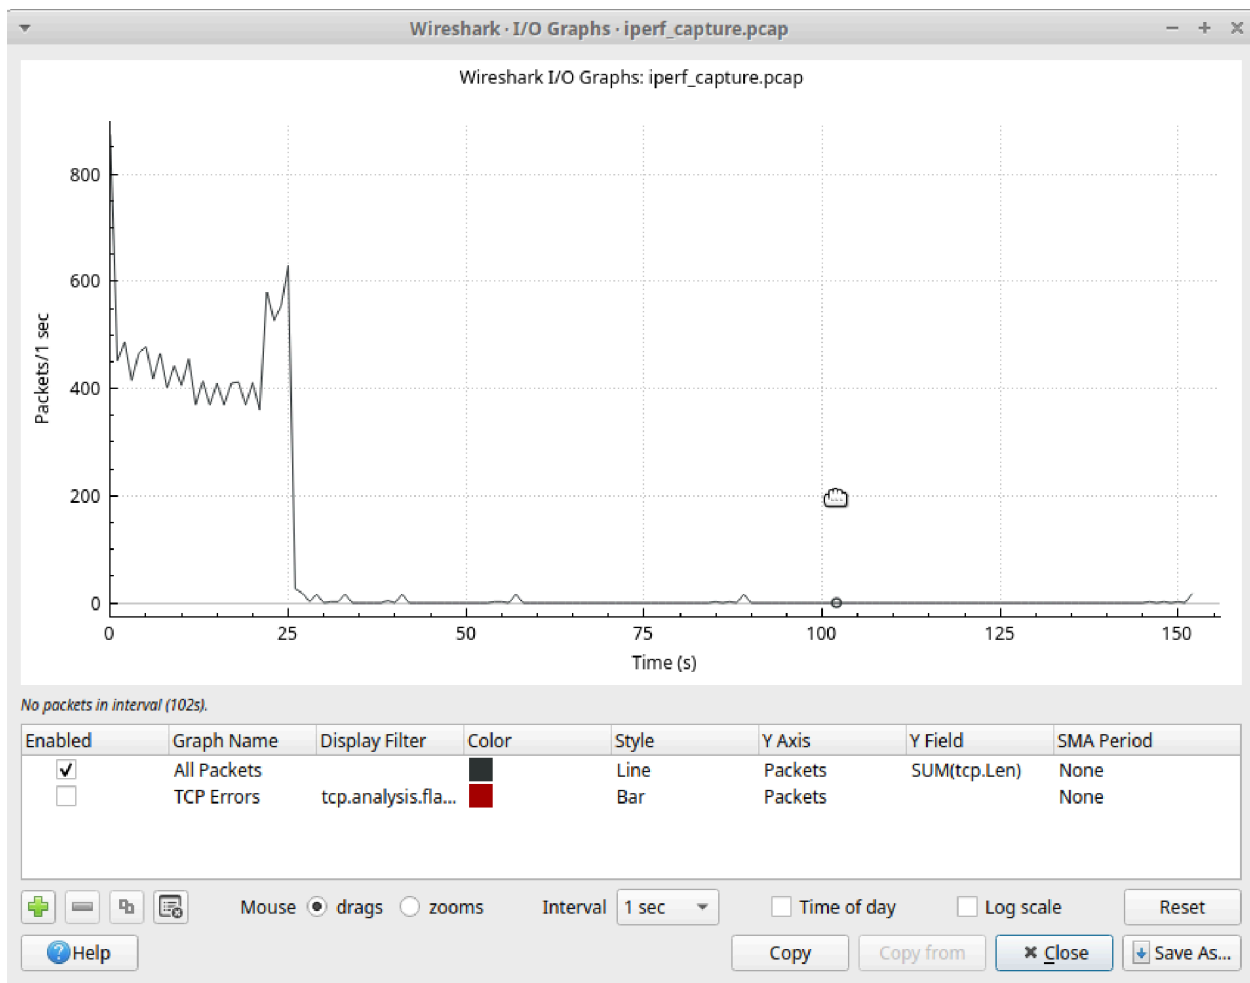
Packet Loss: 1%

- Congestion Control = 'westwood'

Throughput:

Time Interval (s)	Packets	Bytes
0 <> 10	4895	126152928
10 <> 20	3984	124781344
20 <> 30	3124	70349044
30 <> 40	24	1584
40 <> 50	16	1056

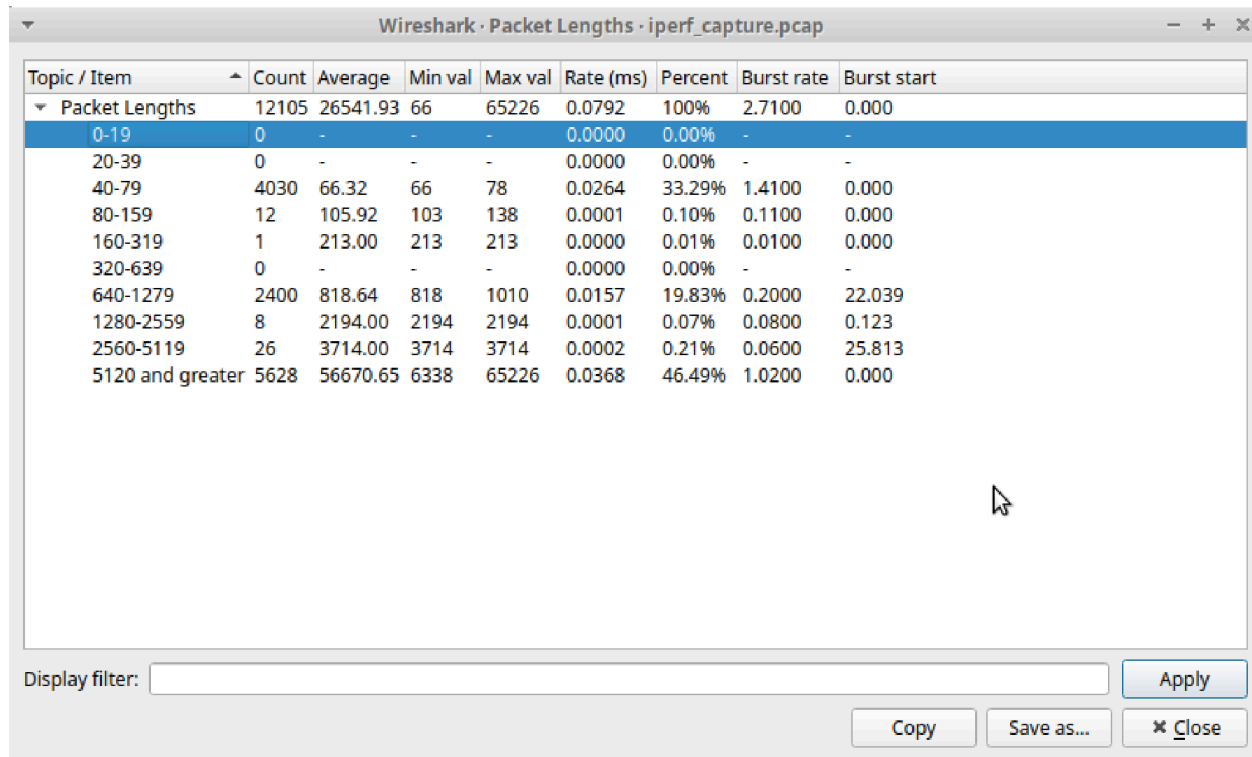
50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	20	1320
90 <> 100	0	0
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	4	264
150 <> Dur	18	1189



Goodput: 69%

Packet Loss: 1

max window size: 61452 bytes



Topic / Item	Count	Average	Min val	Max val	Rate (ms)	Percent	Burst rate	Burst start
▼ Packet Lengths	12105	26541.93	66	65226	0.0792	100%	2.7100	0.000
0-19	0	-	-	-	0.0000	0.00%	-	-
20-39	0	-	-	-	0.0000	0.00%	-	-
40-79	4030	66.32	66	78	0.0264	33.29%	1.4100	0.000
80-159	12	105.92	103	138	0.0001	0.10%	0.1100	0.000
160-319	1	213.00	213	213	0.0000	0.01%	0.0100	0.000
320-639	0	-	-	-	0.0000	0.00%	-	-
640-1279	2400	818.64	818	1010	0.0157	19.83%	0.2000	22.039
1280-2559	8	2194.00	2194	2194	0.0001	0.07%	0.0800	0.123
2560-5119	26	3714.00	3714	3714	0.0002	0.21%	0.0600	25.813
5120 and greater	5628	56670.65	6338	65226	0.0368	46.49%	1.0200	0.000

Display filter:  Apply

Copy Save as... ✖ Close

## Observations

- All the congestion control methods perform very similarly.
- WestWood gives the best throughput because of its conservative congestion window strategy.

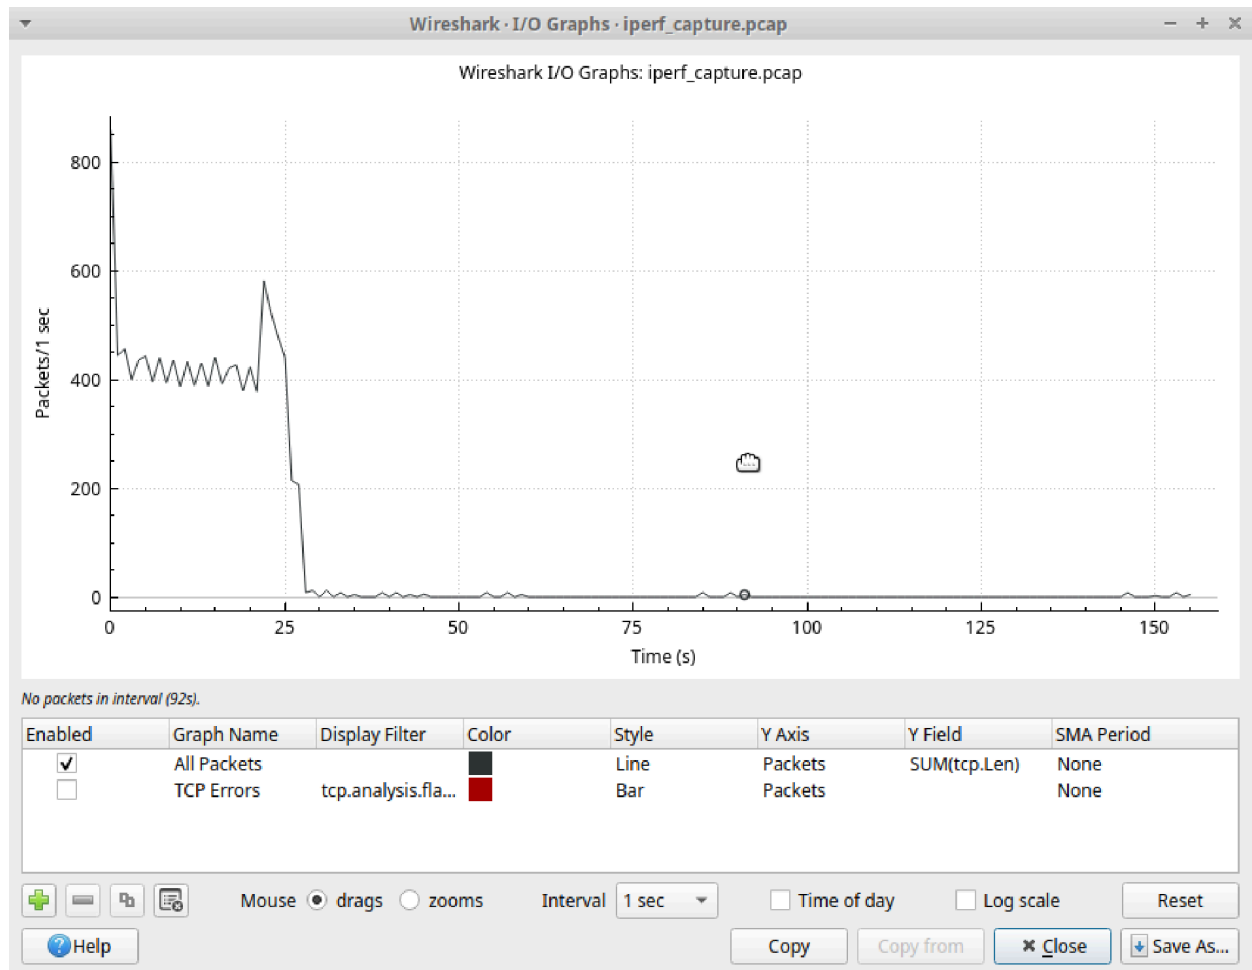
## Part (b)

- Congestion Control = 'yeah'

Throughput:

Time Interval (s)	Packets	Bytes
0 <> 10	4702	126140186
10 <> 20	4087	124788184

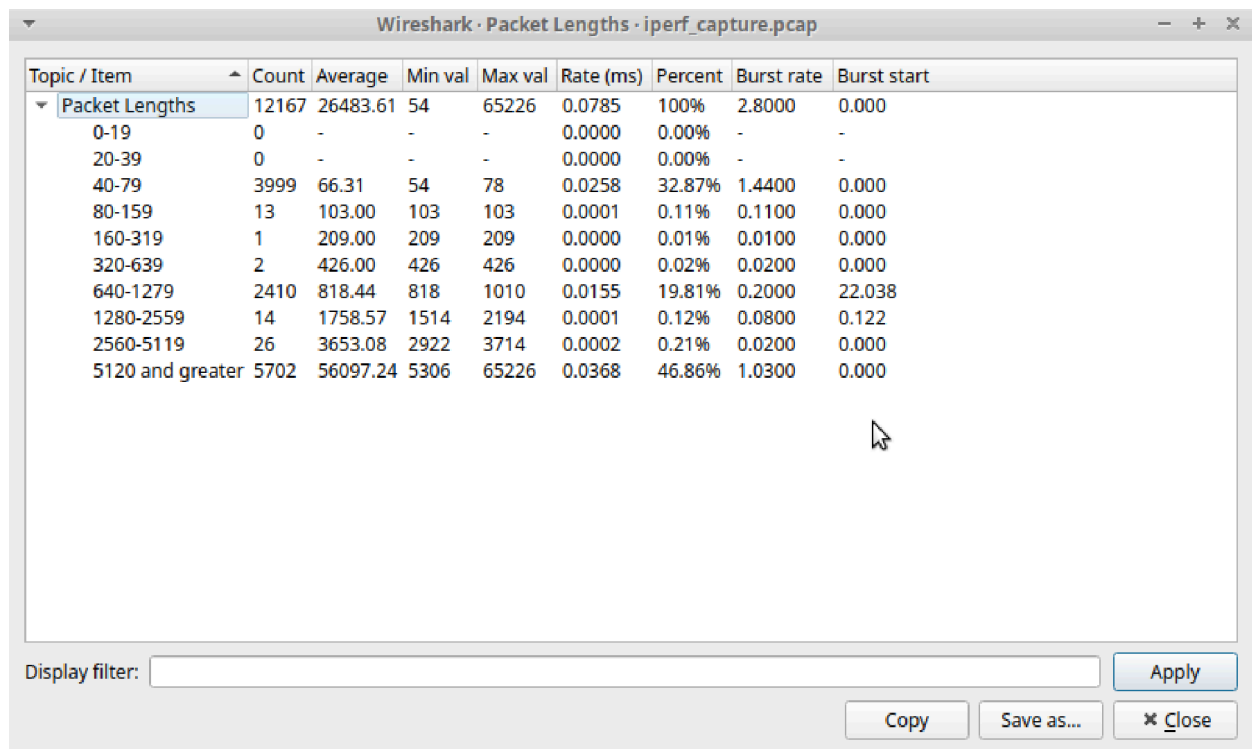
20 <> 30	3267	71290346
30 <> 40	32	2112
40 <> 50	17	1175
50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	16	1056
90 <> 100	4	264
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	8	528
150 <> Dur	14	925



Goodput: 68%

Packet Loss: 1.5

Maximum Packet Size: 65226

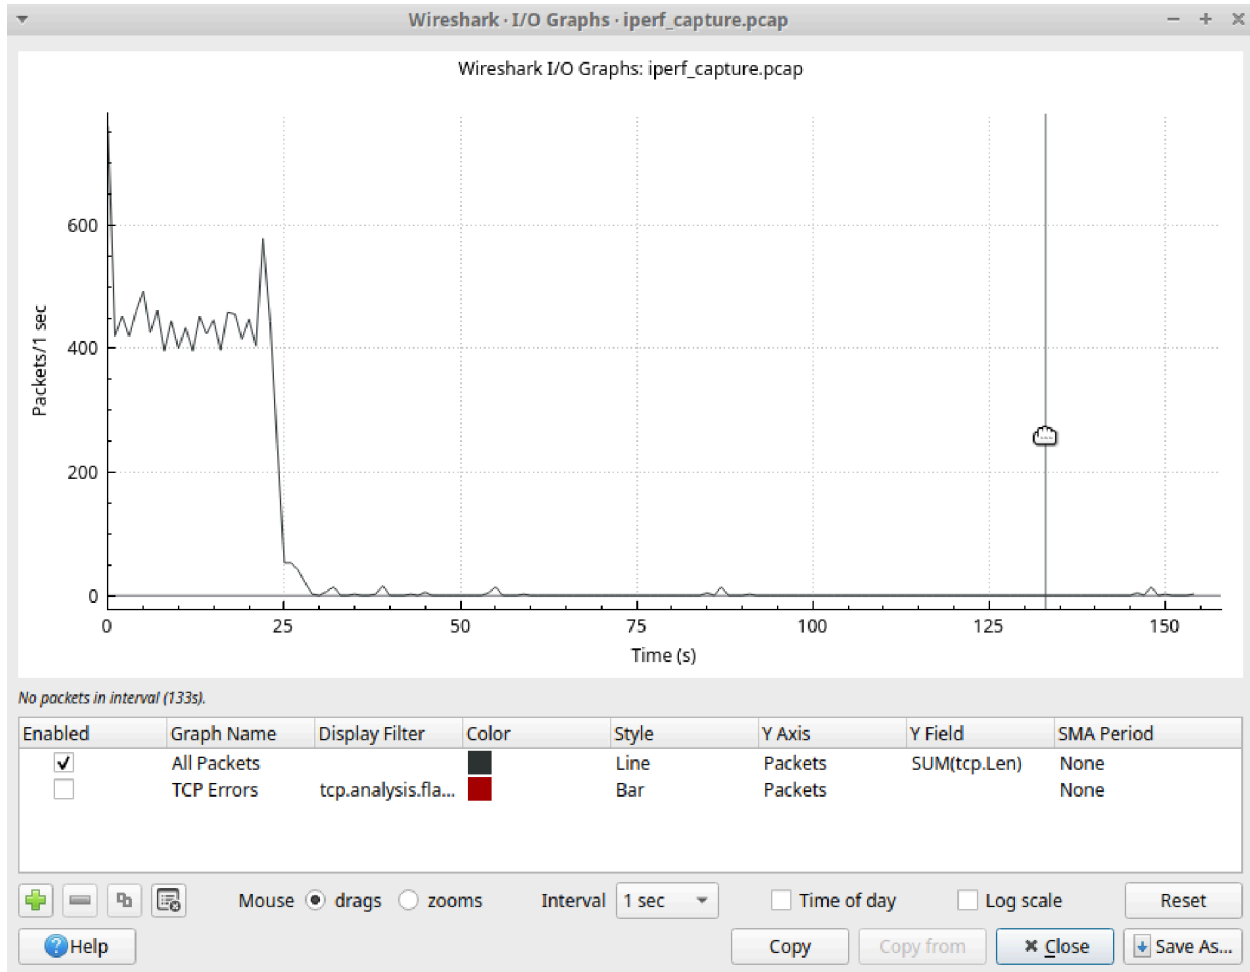


- Congestion Control = 'bbr'

Throughput:

Time Interval (s)	Packets	Bytes
0 <> 10	4733	126142231
10 <> 20	4279	124800868
20 <> 30	2288	60280764
30 <> 40	40	2640
40 <> 50	7	515
50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	18	1188
90 <> 100	2	132
100 <> 110	0	0

110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	18	1188
150 <> Dur	4	265



Goodput: 60%

Packet Loss: 2%

Maximum Packet Size: 614579

Wireshark · Packet Lengths · iperf\_capture.pcap

Topic / Item	Count	Average	Min val	Max val	Rate (ms)	Percent	Burst rate	Burst start
▼ Packet Lengths	11409	27279.44	66	65226	0.0736	100%	2.4100	0.000
0-19	0	-	-	-	0.0000	0.00%	-	-
20-39	0	-	-	-	0.0000	0.00%	-	-
40-79	4161	66.26	66	78	0.0269	36.47%	1.4200	0.000
80-159	13	103.00	103	103	0.0001	0.11%	0.1100	0.000
160-319	7	271.43	208	282	0.0000	0.06%	0.0700	0.000
320-639	0	-	-	-	0.0000	0.00%	-	-
640-1279	2423	818.00	818	818	0.0156	21.24%	0.2000	22.037
1280-2559	0	-	-	-	0.0000	0.00%	-	-
2560-5119	0	-	-	-	0.0000	0.00%	-	-
5120 and greater	4805	64301.80	6210	65226	0.0310	42.12%	0.8100	0.000

Display filter:  Apply

Copy Save as... ✕ Close

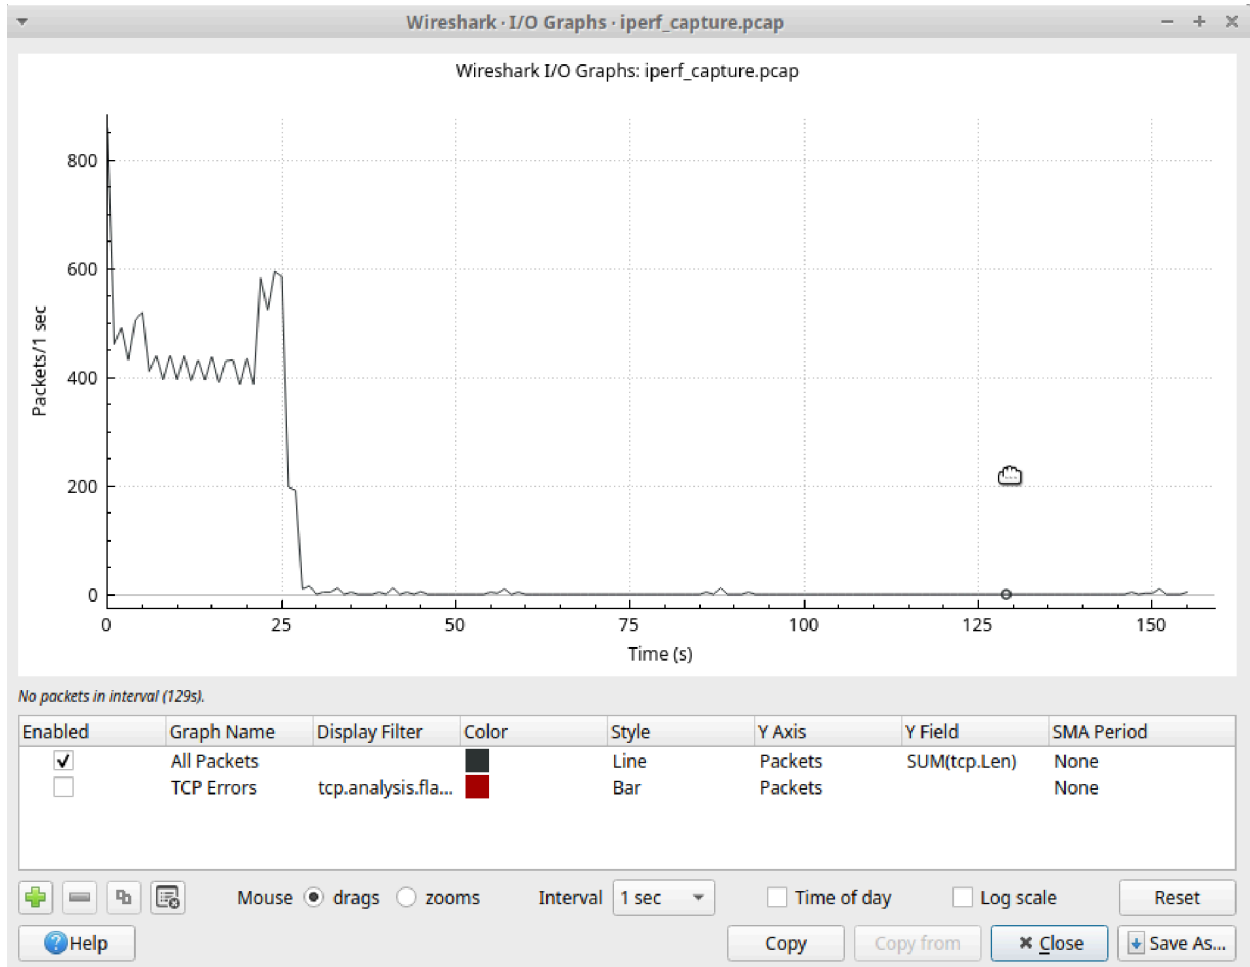
- Congestion Control = 'westwood'

Throughput:

Time Interval (s)	Packets	Bytes
0 <> 10	4955	126163000
10 <> 20	4133	124791232
20 <> 30	3526	75895348
30 <> 40	28	1848
40 <> 50	21	1439
50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	16	1056
90 <> 100	4	264
100 <> 110	0	0



110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	6	396
150 <> Dur	16	1057



Goodput: 65

Packet Loss: 1

## Observations

- YeAH reduces bandwidth competitively so we see more continuous spiky variations


- BBR adjusts based on measured bandwidth and round-trip time, ensuring better stability.
  - Westwood relies on estimated available bandwidth, meaning it takes longer to ramp up in shared environments.
- 

## Part (c)

- Condition 1
  - Congestion Control = 'yeah'

Throughput:

Duration Range	Count	Value
0 <> 10	5029	126161768
10 <> 20	4481	124814146
20 <> 30	3694	77055220
30 <> 40	26	1716
40 <> 50	18	1188
50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	2	132
90 <> 100	18	1188
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	2	132
150 <> Dur	20	1321



Goodput: 60

Packet Loss: 1.5


Mps: 65226

- Congestion Control = 'bbr'

Throughput:

Duration Range	Count	Value
0 <> 10	5029	126161768
10 <> 20	4481	124814146
20 <> 30	3694	77055220
30 <> 40	26	1716
40 <> 50	18	1188

50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	2	132
90 <> 100	18	1188
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	2	132
150 <> Dur	20	1321



Goodput:65


Packet Loss:0.5

MPS: 65226

- Congestion Control = 'westwood'

Throughput:

Duration Range	Count	Value
0 <> 10	5029	126161768
10 <> 20	4481	124814146
20 <> 30	3694	77055220
30 <> 40	26	1716
40 <> 50	18	1188
50 <> 60	20	1320
60 <> 70	0	0
70 <> 80	0	0
80 <> 90	2	132
90 <> 100	18	1188
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	2	132
150 <> Dur	20	1321




Goodput: 68

Packet Loss: 1.5

MPS: 65226

- Condition 2
  - Congestion Control = 'yeah'

Throughput:




Goodput:65

Packet Loss:0.5

MPS: 65226

- Congestion Control = bbr'

Throughput:



Goodput:65


Packet Loss:0.5

MPS: 65226

- Congestion Control = 'westwood'

Throughput:





Goodput:65

Packet Loss:0.5


MPS: 65226

- Congestion Control = 'bbr'

Throughput:

- Condition 3
  - Congestion Control = 'yeah'

Throughput:




Goodput:65

Packet Loss:0.5

MPS: 65226

- Congestion Control = 'bbr'

Throughput:




Goodput:61

Packet Loss:1

- Congestion Control = 'westwood'

Throughput:



Goodput: 67

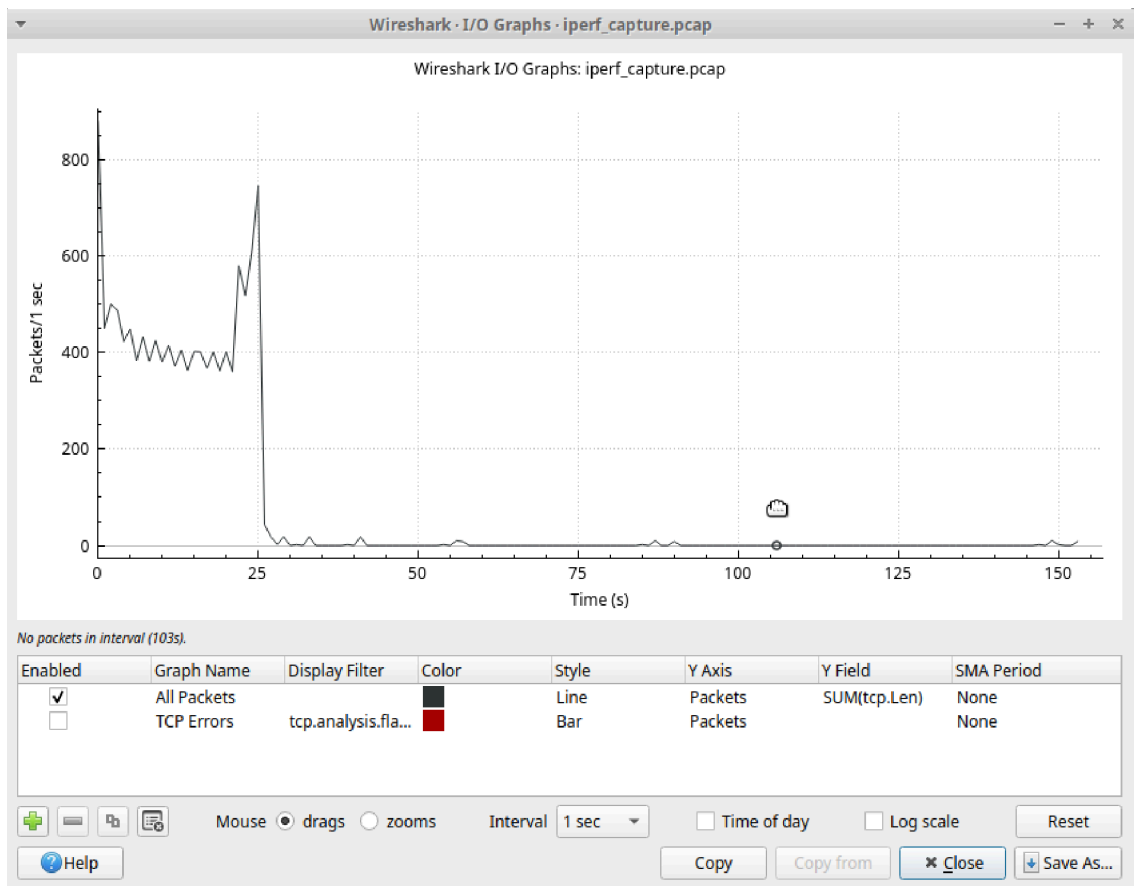
Packet Loss: 0.5

## Observations:

- YeAH is affected as it reacts to loss by switching modes.
- BBR remains steady because it does not use packet loss as a congestion signal.
- Westwood sees degradation since it estimates bandwidth based on successful transmissions.

## Part (d)

- Condition 1: loss = 1%
  - Congestion Control = 'yeah'



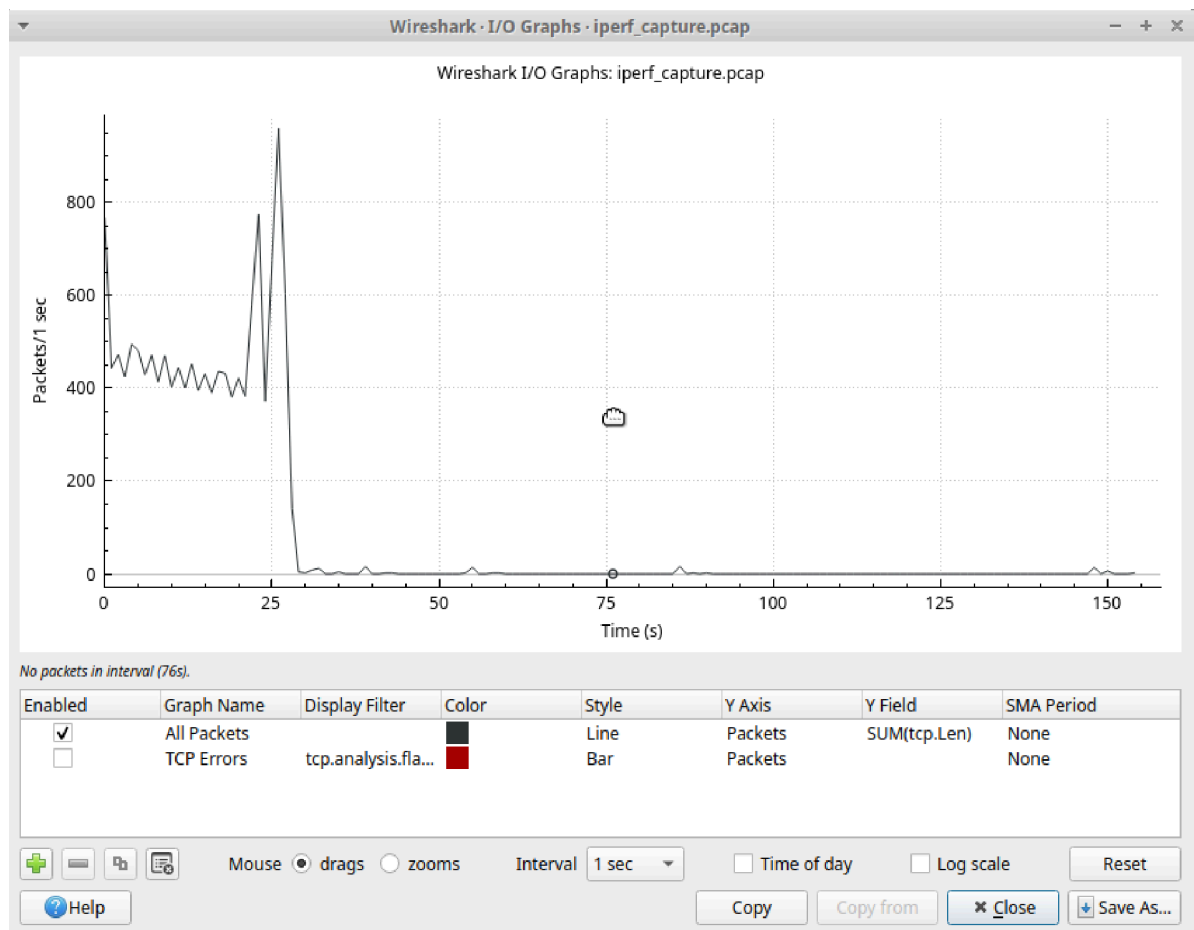
Throughput:

Goodput:

Packet Loss:

- Congestion Control = 'bbr'

Throughput:




Goodput:

Packet Loss:

Maximum Packet Size:

- Congestion Control = 'westwood'

Throughput:



Goodput:

Packet Loss:

Maximum Packet Size:


Condition 2: loss = 5%

- Congestion Control = 'yeah'

Throughput:

	0 <> 10		4958		126157082	
	10 <> 20		4188		124794808	
	20 <> 30		3658		77308764	
	30 <> 40		26		1716	
	40 <> 50		18		1188	
	50 <> 60		20		1320	
	60 <> 70		0		0	

70 <> 80	0	0
80 <> 90	2	132
90 <> 100	18	1188
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	2	132
150 <> Dur	20	1321



Goodput:

Packet Loss:


Maximum Packet Size:

- Congestion Control = 'bbr'



Throughput:

0 <> 10	5076	126164869
10 <> 20	4415	124809790
20 <> 30	6342	84759756
30 <> 40	160	3523904
40 <> 50	12	792
50 <> 60	16	1056
60 <> 70	4	264
70 <> 80	0	0
80 <> 90	10	660
90 <> 100	10	660
100 <> 110	0	0
110 <> 120	0	0
120 <> 130	0	0
130 <> 140	0	0
140 <> 150	8	528
150 <> Dur	14	925




Goodput:

Packet Loss:

Maximum Packet Size:

- Congestion Control = 'westwood'

Throughput:



Goodput:

Packet Loss:

Maximum Packet Size:

## Observations

- YeAH is the best performer because it ignores loss-based congestion signals and estimates available bandwidth.
- BBR suffers more as loss increases because it switches between fast and slow modes.
- Westwood performs worst under high loss because it relies on bandwidth estimation, which gets inaccurate with lost ACKs.

# Task 2: Implementation and mitigation of SYN flood attack

## Setting up the environment for attack

### 1. Server

- a. disable cookies
- b. increase the number of retries
- c. decrease the size of backlog so that legitimate traffic is lost

```
sudo sysctl -w net.ipv4.tcp_syncookies=0
sudo sysctl -w net.ipv4.tcp_synack_retries=255
sudo sysctl -w net.ipv4.tcp_max_syn_backlog=16
```

### 2. Client

- a. Starting packet capture and legitimate traffic

```
sudo tcpdump -i wlp0s20f3 host 172.20.10.2 -w Q2-part1.pcap > out.txt
while true; do curl http://172.20.10.2:8000/;sleep 1;done
```

- b. Starting attack

```
sudo iptables -A OUTPUT -p tcp --tcp-flags RST RST -j DROP
sudo hping3 -S -p 8000 -i u10000 172.20.10.2
```

- -S means SYN flag
- -i means send 1 packet in 10000 micro sec = 100 packets per second

- c. Stopping attack

```
sudo iptables -D OUTPUT -p tcp --tcp-flags RST RST -j DROP
```

- **ctrl+c** to stop the terminal running **hping3**

## For mitigation

### 1. Server

#### a. Restoring the default settings

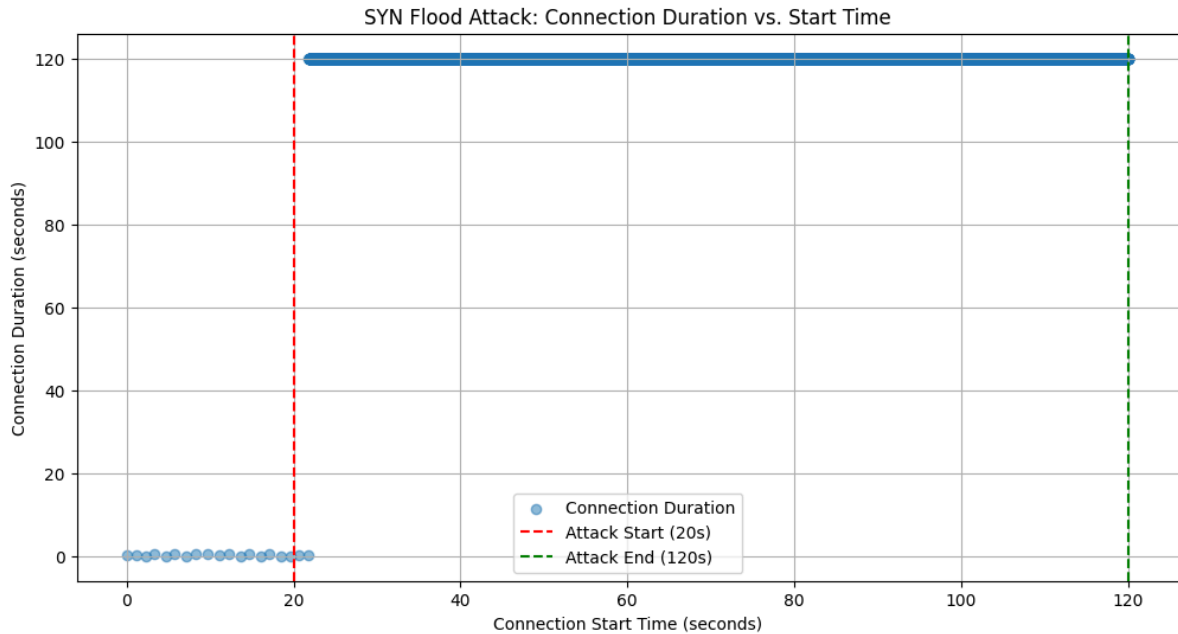
```
sudo sysctl -w net.ipv4.tcp_syncookies=0  
sudo sysctl -w net.ipv4.tcp_synack_retries=255  
sudo sysctl -w net.ipv4.tcp_max_syn_backlog=16
```

```
sudo sysctl -w net.ipv4.tcp_syncookies=1  
sudo sysctl -w net.ipv4.tcp_synack_retries=5  
sudo sysctl -w net.ipv4.tcp_max_syn_backlog=128
```

### 2. Client remains same


## Results

### 1. Successful attack



- Very low connection duration in 0 - 20 seconds (legitimate traffic)
- Unclosed connections in 20 - 120 seconds (attack period)
- No connections accepted after 120 seconds because server is busy retrying the old attack packets
- Legitimate traffic stopped getting response after first 20 seconds. (DoS)

## 2. Mitigated attack



- The legitimate traffic is maintained throughout
- The connections opened in the attack interval are never close and stay at duration = 100 (default)
- Utilize cookies and bigger backlog buffer to store incoming requests and reduce the retries.

## Task 3: Analyze the effect of Nagle's algorithm on TCP/IP performance

### Setup and Tools Used

- UTM Virtual Machine Ubuntu 20.04
- Mininet
- Wireshark

### Implementation

#### Setting Up Network

```

def setup_topology():

    class CustomTopo(Topo):
        def build(self):
            s1 = self.addSwitch('s1')
            h1 = self.addHost('h1')
            h7 = self.addHost('h7')

            self.addLink(h1, s1)
            self.addLink(s1, h7)

    net = Mininet(topo=CustomTopo(), controller=OVSController)
    net.start()

    net.get('s1').cmd('ovs-ofctl add-flow s1 actions=normal')
    return net


def configure_tcp_options(host, nagle, delayed_ack):

    if nagle:
        host.cmd("sysctl -w net.ipv4.tcp_nodelay=0")
    else:
        host.cmd("sysctl -w net.ipv4.tcp_nodelay=1")

    if delayed_ack:
        host.cmd("sysctl -w net.ipv4.tcp_delack_min=100")
    else:
        host.cmd("sysctl -w net.ipv4.tcp_delack_min=0")

```

## Running Test

```

def run_tcp_test(net, nagle, delayed_ack):

```



```

server = net.get('h7')
client = net.get('h1')

configure_tcp_options(server, nagle, delayed_ack)
configure_tcp_options(client, nagle, delayed_ack)

# Get correct IP for h7
server_ip = server.IP()

# Start tcpdump on the correct interface
server.cmd('tcpdump -i h7-eth0 port 5001 -w tcp_test.pcap &')
time.sleep(3) # Ensure tcpdump starts before traffic

# Start iPerf Server
server.cmd('iperf -s -p 5001 &')
time.sleep(2) # Ensure server starts

# Start iPerf Client
client.cmd(f'iperf -c {server_ip} -p 5001 -t 120 -b 320bps')

time.sleep(35) # Ensure enough time for data transfer

server.cmd('pkill iperf')
time.sleep(2)
server.cmd('pkill tcpdump')
time.sleep(2)

print("Test completed! Results saved to tcp_test.pcap.")
analyze_pcap()

```

## Analysing

```

def analyze_pcap():

    if not os.path.exists("tcp_test.pcap") or os.path.getsize("tcp_test.pcap") == 0:

```

```
print("[ERROR] No packets were captured. Check tcpdump settings!")
return
```

```
throughput = subprocess.getoutput("tshark -r tcp_test.pcap -q -z io,stat,60 | c
goodput = subprocess.getoutput("tshark -r tcp_test.pcap -Y 'tcp.len > 0' | wc
packet_loss = subprocess.getoutput("tshark -r tcp_test.pcap -q -z expert,ip | c
max_packet_size = subprocess.getoutput("tshark -r tcp_test.pcap -T fields -e
```

```
print("Throughput:", throughput)
print("Goodput (total data packets received):", goodput)
print("Packet loss rate:", packet_loss)
print("Maximum packet size achieved:", max_packet_size)
```

## Results

### Part (a) Nagle Disabled, Delayed ACK Disabled

Throughput:

Time Stamp	Packets	Total bytes transferred
0.000 <> 0.032	8	4640

Goodput: 16.7

Packet Loss Rate: 0

Maximum Packet Size: 682

### Part (b) Nagle Disabled, Delayed ACK Enabled

Throughput:

0.000 <> 0.012	8	4640
----------------	---	------

Goodput: 12.5

Packet Loss Rate: 0

### Part (c) Nagle Enabled, Delayed ACK Disabled

Throughput:

| 0.000 <> 0.035 | 8 | 4640 |

Goodput: 12.5

Packet Loss Rate: 0

## Part (d) Nagle Enabled, Delayed ACK Enabled

Throughput:

| 0.000 <> 0.032 | 8 | 4640 |

Goodput: 11.5

Packet Loss Rate: 0

Maximum Packet Size: 768

## Explanation

### 1. Nagle Enabled, Delayed ACK Enabled

- **Nagle buffers packets** instead of sending immediately.
- **Delayed ACKs wait** before acknowledging, causing additional delays.
- **Result:** Very inefficient transfer, **low throughput and goodput**.

### 2. Nagle Enabled, Delayed ACK Disabled

- **Nagle still buffers**, but ACKs come faster since DA is off.
- **Result:** Slightly better performance but still **moderate delays**.

### 3. Nagle Disabled, Delayed ACK Enabled

- **Packets are sent immediately**, but ACKs are delayed.
- **Result:** More packets in flight, but **some unnecessary retransmissions** due to delayed ACKs.

### 4. Nagle Disabled, Delayed ACK Disabled

- **Every packet is sent immediately** (no waiting for Nagle).
- **ACKs are sent immediately** (no waiting for DA).

- **Result:** Best efficiency, **highest throughput and goodput.**

## Observations

- IMP: Because TCP is a dynamic protocol, even when we limit the bandwidth of TCP connection, the backend utilises the maximum possible bandwidth.
- We observe the outcomes of the execution similar to explanation above.
- Part (a) achieves highest goodput because of lower delays and there is no need for conservative congestion control.

## Acknowledgement

We would like to thank Prof. Sameer Kulkarni to provide this opportunity to work on this assignment. We also thank teaching assistants for their constant support.